



UNIVERSITAT POLITÈCNICA DE CATALUNYA
BARCELONATECH

Escola Tècnica Superior d'Enginyeria
de Telecomunicació de Barcelona



Deep Learning for Neurodegenerative Disease detection through Spontaneous Speech

Master Thesis

submitted to the Faculty of the
Escola Tècnica d'Enginyeria de Telecomunicació de Barcelona
Universitat Politècnica de Catalunya

by

Roger Esteve

In partial fulfillment
of the requirements for the master in
*Master's degree in Advanced Telecommunication Technologies (MATT), track in Deep
Learning for Multimedia Processing*

Advisor: Dr Francisco Javier Hernando Pericas
Barcelona, June 2026



Contents

List of Figures	4
List of Tables	4
1 Introduction	7
1.1 Motivation and objectives	8
1.2 Work plan	9
2 State of the art of the technology used or applied in this thesis	10
2.1 Biological motivation: amyloid pathology and clinically overlapping syn- dromes	10
2.2 Spoken language as a scalable digital biomarker	11
2.3 From shallow classifiers to deep neural representations	11
2.3.1 Self-supervised acoustic encoders: Wav2Vec 2.0	11
2.3.2 Contextual text encoders	12
2.4 Benchmark challenges and public corpora	13
2.5 Automatic differentiation of PPA variants from connected speech	13
2.6 Multimodal fusion: pooling, cross-attention, and temporal alignment	14
2.6.1 Self-attention fusion	14
2.6.2 Cross-attention and gated variants	15
2.6.3 Selective state-space models: Mamba	15
2.6.4 Cross-modal alignment	15
2.7 Handcrafted acoustic descriptors alongside learned embeddings	17
2.8 Synthesis and positioning	17
3 Methodology and Project Development	18
3.1 Problem Formulation	18
3.2 Datasets and Preprocessing	19
3.2.1 ADReSSo Corpus	19
3.2.2 SPIN Clinical Cohort	20
3.2.3 ASR Pipeline and Transcript Quality	21
3.3 System Architecture	21
3.3.1 Acoustic Encoder: Choice and Configuration	22
3.3.2 Wav2Vec Pooling and Layer Mixing	22
3.3.3 Text Encoder: Choice and Configuration	24
3.3.4 Feature-Extractor Stacks	24
3.3.5 Cross-Modal Alignment	25
3.3.6 Fusion Module	25
3.3.7 Handcrafted Feature Extraction and Integration	28
3.3.8 Classification Heads	30
3.4 Unified Implementation: <code>res_model</code>	30
3.5 Training Protocol	31
3.6 Ablation Protocol	31
3.7 Evaluation Metrics and Baselines	32

4	Results	33
4.1	Alzheimer’s Disease Classification (ADReSSo)	33
4.2	Amyloid- β Status Prediction (SPIN)	34
4.3	PPA Variant Classification (WAB)	34
4.4	Ablation Studies	34
4.4.1	Modality	35
4.4.2	Handcrafted Acoustic Features	35
4.4.3	Fusion	35
4.4.4	Wav2Vec Layer Pooling and the Layer Mixer	36
4.4.5	Audio–Text Alignment	38
4.4.6	Feature Extractors	40
4.4.7	Unified-Model Confirmation	40
4.5	Error Analysis	40
5	Budget	42
6	Environment Impact (Optional)	43
7	Conclusions and Future Development	44
	References	45
	Appendices	49

List of Figures

1	Gantt diagram of the project.	9
2	Architecture of the Wav2Vec 2.0 self-supervised speech encoder.	12
3	Overview of the three multimodal fusion paradigms: early, intermediate, and late.	14
4	Cross-modal alignment between acoustic frames and text tokens, comparing unsynchronised input (top) with word-to-token aligned input (bottom). . .	16
5	End-to-end architecture used in this thesis: frozen Wav2Vec 2.0 and BERT-family encoders produce acoustic and textual sequences, which are placed on a common index by the alignment step, combined by the sequence-to-sequence fusion module, and mapped to a task label by a lightweight head. . .	23
6	Audio-text alignment schemes compared in this thesis: the CogniAlign word-to-token baseline, and the two proposed token-to-token schemes (uniform, which spreads each word's frames evenly across its subword tokens, and proportional, which allocates frames to tokens in proportion to their character length).	26
7	Self-attention fusion over the concatenated audio-text sequence.	27
8	Cross-attention fusion variants: (a) standard, (b) gated, and (c) bidirectional. . .	28
9	Mamba selective state-space fusion applied to the concatenated audio-text sequence.	29
10	Validation accuracy of the fusion families on the three tasks (values from Table 8). The cross-attention family leads on every task, while Mamba's drop to chance on PPA stands out as the clearest outlier.	37
11	Validation accuracy against per-sample inference time for the fusion families (Step B, ADReSSo). Cross-attention pairs high accuracy with low cost in the precomputed-embedding pipeline; absolute timings depend on whether the encoders are cached or run in the loop.	37
12	Learned per-layer weights of the per-utterance sample mixer on ADReSSo ($\tau = 0.1$, twelve Wav2Vec 2.0 layers). Weight concentrates on the uppermost layer with secondary mass on the lowest, and the per-utterance spread is shown as a shaded band around the mean.	39

Listings

List of Tables

1	Clinical cohorts and evaluation protocol	19
2	SPIN demographics by amyloid status	21
3	SPIN demographics by PPA variant	21
4	Fusion families	27
5	Handcrafted acoustic features	29
6	Modality ablation per task	35
7	Handcrafted acoustic-feature ablation per task	36

8	Fusion-technique ablation per task	38
9	Wav2Vec layer-pooling ablation per task	38
10	Audio-text alignment ablation per task	39
11	Feature-extractor single-factor ablation per task	40
12	Best technique bundle per task	41

Abstract

Speech is an attractive non-invasive biomarker for neurodegenerative disease, but multi-modal systems that combine acoustic and textual signals expose many interacting design choices whose individual contributions are rarely isolated. This thesis treats such a system as a composition of largely independent design axes (modality, fusion, audio-text alignment, Wav2Vec layer pooling, layer mixing, and feature extractor) and estimates the contribution of each through controlled ablations across three clinical tasks: Alzheimer’s disease classification on ADReSSo, amyloid- β status prediction, and primary progressive aphasia variant classification. Its central methodological contribution is a pair of token-to-token alignment schemes, uniform and proportional, that refine the word-to-token alignment of prior work at the subword level. The strongest per-axis choices are consolidated into a single configuration-driven implementation, yielding a recommended architecture that is competitive across tasks while remaining lightweight and deployable.

1 Introduction

Alzheimer’s disease is the most common cause of dementia worldwide. Its underlying pathology, the progressive build-up of amyloid- β plaques and neurofibrillary tau tangles, begins years before any clinical symptom appears [1], which is why detecting it early has become a central goal of current research. Confirmation is the bottleneck: a definitive read on amyloid status still depends on a lumbar puncture or a PET scan, neither of which scales to routine screening. Diagnosis at the syndrome level is no easier. Primary progressive aphasia (PPA) comes in three canonical variants, non-fluent/agrammatic, semantic, and logopenic [2, 3], whose behavioural profiles overlap enough that telling them apart is difficult, especially in early or atypical cases [4]. Scalable amyloid screening and reliable PPA subtyping thus point to the same need: non-invasive computational tools that support, rather than replace, expert judgement.

Speech is a natural place to look. Producing a spontaneous description draws at once on cognitive planning, lexical access, syntactic construction, motor execution, and prosodic control, a chain of processes that neurodegeneration disrupts unevenly and so leaves its traces in the signal. And unlike a scan or an assay, a recording costs almost nothing to collect and can be gathered repeatedly, even remotely [5]. The clinical infrastructure is already in place: structured prompts such as the Picnic Scene from the Western Aphasia Battery [6] are routine in neuropsychological batteries, and the recordings they yield can be analysed computationally with no extra instrumentation.

The first computational attempts leaned on handcrafted linguistic descriptors and a few shallow acoustic statistics; they proved the idea workable but generalised poorly [7]. Deep learning raised the ceiling. Self-supervised encoders, Wav2Vec 2.0 on the acoustic side [8] and BERT-family models on the textual [9], supply far richer representations than any hand-built feature set. Shared benchmarks made the gains measurable: the ADReSSo challenge [10] fixed splits and protocols for like-for-like comparison, and the surveys built on this line of work converge on one finding in particular, that combining acoustic and textual evidence beats either modality alone [11].

Yet a headline accuracy number hides as much as it reveals. A typical pipeline bundles several decisions at once, which modalities to use, how to fuse them, how to align acoustic frames with text tokens, how to configure the self-supervised encoder, and then reports a single end-to-end figure. With the choices entangled inside one system, there is no way to tell which technique earned the performance and which merely came along for the ride on a particular dataset or implementation. A second blind spot is linguistic. Most benchmarked systems are built and tested on English, and whether their trends survive in languages with different morphology and prosody is largely unknown. Pathology-confirmed cohorts in other languages do exist—the Spanish-speaking Sant Pau Initiative on Neurodegeneration (SPIN) [12] among them—but they have drawn little attention in the multimodal deep-learning literature.

1.1 Motivation and objectives

These two gaps, entangled design choices and an English-centric evidence base, set the agenda for this thesis. The premise is straightforward: if published systems differ along every axis at once, the only way to learn what each axis contributes is to move one at a time. A structured ablation programme, varying a single design choice while the rest stay pinned to controlled defaults, turns a tangle of confounded comparisons into per-axis evidence that can guide later design.

An earlier companion study [13] reported first results for AD and PPA classification on the ADReSSo and SPIN cohorts; the present thesis builds on it by breaking the pipeline into design choices that can each be ablated in isolation. Its main objectives are:

1. **Modality.** Quantify the individual and combined contribution of audio, text, and handcrafted acoustic features across three clinical speech tasks: AD classification (ADReSSo, English), amyloid- β status prediction (SPIN, Spanish, validation-only), and PPA variant classification (the same SPIN cohort, WAB picture-description subset, Spanish, validation-only).
2. **Fusion.** Systematically compare seven sequence-to-sequence fusion families (self-attention, multi-head self-attention, transformer with positional encoding, cross-attention, gated cross-attention, bidirectional cross-attention, and a Mamba state-space comparator), on all three tasks, taking bidirectional cross-attention as the CogniAlign default and replicating the trends under a second, independent implementation.
3. **Alignment.** Propose and evaluate two novel token-to-token audio-text alignment schemes (uniform and proportional) that operate at the subword level, benchmarked against the word-to-token alignment of CogniAlign [14]; this is the central methodological contribution of the thesis.
4. **Wav2Vec pooling.** Evaluate final-layer versus weighted multi-layer Wav2Vec pooling, and an adaptive per-utterance *sample* layer mixer evaluated in the unified implementation.
5. **Feature extractor.** Conduct a single-factor encoder ablation that swaps one audio or text backbone at a time while holding the rest of the pipeline fixed, spanning 24 runs (four encoder variants $\times$ three tasks $\times$ two seeds).
6. **Recommended architecture.** Synthesise the strongest per-axis choices into a recommended hybrid architecture, developed in the Discussion.

The programme works through the axes in stages: a modality study (Step A), a fusion-family comparison (Step B), an alignment study that sets the three schemes against one another under a fixed fusion at the final Wav2Vec layer (Tier A), a pooling study weighing the final layer against a learnable weighted sum (WvLayer), and a combined study crossing alignment with weighted pooling (Align $\times$ Wv). The strongest choice from each stage is then assembled and re-validated in a single configuration-driven implementation, **res_model**. A second stack, SER, plays only a supporting role: it has no alignment axis and is not a rival to be beaten, but an independent codebase on which the fusion and

pooling trends can be checked for reproduction. What the thesis compares throughout is the technique, not the software that happens to implement it.

The remainder of this thesis is organised as follows. Section 2 reviews the state of the art at the intersection of neurodegenerative-disease biomarkers, clinical linguistics, and multimodal deep learning on spontaneous speech. Section 3 formalises the prediction problem as a composition of six design axes, describes the three clinical cohorts and their preprocessing, the system architecture and its three implementations, the training protocol, and the evaluation metrics. Section 4 reports the staged ablation programme and the unified-model validation, and Section 7 synthesises the findings into the recommended architecture and outlines future work.

1.2 Work plan




Figure 1: Gantt diagram of the project.

2 State of the art of the technology used or applied in this thesis

This chapter surveys relevant advances at the intersection of neurodegenerative disease biomarkers, clinical linguistics, and multimodal deep learning on spontaneous speech.

This chapter surveys the scientific and technical background that motivates analysis of spontaneous speech with deep multimodal models in neurodegenerative disease. The discussion follows the structure of recent surveys and challenge-led benchmarking in speech-based cognitive assessment [11], but it is organised around the specific threads that underpin the companion contribution developed in this thesis [13]: pathology-aware biomarker estimation from voice, differential diagnosis of primary progressive aphasia (PPA), cross-lingual dataset gaps, self-supervised acoustic encoders paired with language-model text encoders, fusion architectures that align modalities in time, and the complementary role of handcrafted paralinguistic descriptors. Throughout, emphasis is placed on results and modelling trends that are reproducible on public benchmarks, notably the Alzheimer’s Dementia Recognition through Spontaneous Speech Only (ADRESSo) corpus associated with the challenge formulation of Luz et al. [10], and on how those trends inform deployment on clinically curated cohorts such as the Sant Pau Initiative on Neurodegeneration (SPIN) [12].

2.1 Biological motivation: amyloid pathology and clinically overlapping syndromes

Accumulation of amyloid- β ($A\beta$) is a core pathophysiological substrate in Alzheimer’s disease (AD) research and underpins current therapeutic strategies that lean on early biological identification of at-risk individuals [1]. Confirmatory assessment traditionally relies on cerebrospinal fluid assays or molecular neuroimaging; although accurate, these modalities are expensive, invasive, or capacity constrained, which contributes to delayed detection relative to symptom onset and limits scalable screening in population settings.

Diagnosis at the syndrome level is a separate problem, and an equally stubborn one. Language and executive deficits arise across many neurodegenerative aetiologies whose phenotypes only partly overlap, so a given clinical picture seldom maps cleanly onto a single cause. PPA exemplifies this tension: it is defined as a progressive disorder of language as the principal complaint and functional limitation, and contemporary consensus schemes stratify variants according to the dominant cognitive-linguistic profile [2, 3]. The non-fluent/agrammatic (*nfv*), semantic (*sv*), and logopenic (*lv*) variants differ in fluency, grammatical competence, lexical-semantic knowledge, and repetition/spelling-span signatures; importantly, variant labels provide clinically actionable hypotheses about underlying proteinopathy patterns even though symptoms can blur across boundaries [4]. For a model, the consequence is that the discriminating cues are split across two channels: the fine temporal micro-structure of speech (planning pauses, articulatory slowing, reduced informational density) and lexical-semantic content. A system has to read both at once. Both, conveniently, are recoverable from the spontaneous picture descriptions and narrative tasks already collected during routine assessment [15].

2.2 Spoken language as a scalable digital biomarker

Prospective clinicopathological studies and controlled biomarker programmes increasingly incorporate speech and language markers alongside imaging and fluid assays because voice can be acquired frequently, remotely, and at comparatively low marginal cost [5]. At the same time, evaluation methodology matters: speech measures can track syndrome severity and risk enrichment, yet population heterogeneity (education, dialect, comorbid hearing loss, task instructions) can inflate apparent effects unless modelling assumptions are carefully matched to sampling designs [15].

Early computational work emphasised engineered linguistic descriptors extracted from transcripts (lexical diversity, syntactic complexity, idea density) and shallow acoustic statistics computed from short recordings of spontaneous speech to separate AD from healthy ageing [7]. These approaches demonstrated feasibility but often relied on manual transcription pipelines or handcrafted feature dictionaries whose coverage is incomplete for noisy clinical audio. More recent work revisits classical baselines (random forests, support-vector machines) with richer descriptors, but reports sensitivity that swings with task elicitation and recording conditions [15]. The common thread is that these marker families do carry predictive signal, yet shallow decision boundaries leave the temporal dynamics of the waveform largely untapped [16, 17, 18].

2.3 From shallow classifiers to deep neural representations

Representation learning replaces brittle front-ends with hierarchical embeddings trained either supervised end-to-end or semi-supervised with objectives aligned to phonetic/lexical structure.

2.3.1 Self-supervised acoustic encoders: Wav2Vec 2.0

In speech processing broadly, self-supervised pre-training of transformer acoustic models learns frame-level features transferable across languages and domains with comparatively modest labelled downstream data. The Wav2Vec 2.0 architecture [8] illustrates this paradigm. A multi-layer convolutional feature encoder first transforms raw waveform samples into a sequence of latent speech representations at approximately 20-millisecond resolution. These latent vectors are then contextualised by a transformer encoder that attends across the full temporal extent of the utterance. During pre-training, a proportion of the latent frames is masked, and the model solves a contrastive task: for each masked position, it must identify the true quantised representation among distractors sampled from the same utterance. Because no transcriptions or labels are required, the procedure scales to vast quantities of unlabelled audio. The resulting transformer layers learn representations at different levels of abstraction: earlier layers tend to encode acoustic and phonetic detail, while later layers capture more abstract prosodic and linguistic patterns. This layered organisation has practical implications for downstream tasks, which may benefit from aggregating information across multiple transformer layers rather than relying on the final output alone, a question addressed experimentally in this thesis. Two configurations of this model are relevant to the present work. The English **base-960h** model, trained on read English speech, is one; the multilingual **XLSR-300M** variant, which ex-

tends the same framework to around 436,000 hours of speech spanning 128 languages, is the other, and its cross-lingual pre-training makes it better suited to scenarios where labelled pathological speech is scarce in any individual language, such as the Spanish cohorts studied here. The thesis employs both as frozen feature extractors and compares them directly in a controlled encoder ablation, since the more powerful multilingual model is not guaranteed to win once the downstream task is fixed. Related self-supervised models, notably HuBERT [19] and WavLM [20], pursue analogous representation objectives through masked-prediction and denoising tasks respectively, but are not employed in this thesis.




Figure 2: Architecture of the Wav2Vec 2.0 self-supervised speech encoder.

2.3.2 Contextual text encoders

On the text side, the BERT family of encoder-only transformer models [9] learns bidirectional contextual embeddings through masked language modelling: random tokens in the input are masked, and the model is trained to predict them from surrounding context. The resulting representations capture lexical semantics and local syntactic structure, going beyond the static word vectors of earlier approaches. Lighter distilled variants such as DistilBERT [21] preserve much of this capacity in a smaller architecture, while modern encoder-only models such as ModernBERT [22] improve representation quality through updated architectural choices. When automatic speech recognition transcripts are available, these text encoders provide a complementary linguistic view that acoustic features alone may encode only implicitly.

Reviews oriented specifically at AD detection from speech highlight two robust tendencies [11]: deep neural pipelines dominate leaderboard-style benchmarks when adequate

training partitions exist, and multimodal systems that consume both acoustic embeddings and text outperform either modality alone when transcripts are informative enough for alignment. These observations motivate architectures that treat spontaneous speech as *tightly coupled* audio–text sequences rather than as isolated vectors [23].

2.4 Benchmark challenges and public corpora

Community challenges accelerated reproducible comparison by freezing splits and evaluation protocols for cognitive impairment detection from speech-only submissions [10]. The ADRESSo subset derived from the DementiaBank Pitt corpus of picture descriptions remains the de facto English benchmark for Alzheimer’s dementia versus healthy control discrimination under matched recording conditions; numerous challenge systems established competitive baselines [24], and subsequent work continues to extend feature attribution and model explanations [25].

Performance on English benchmarks does not carry over to other languages for free. Morphology, prosodic norms, and clinician elicitation protocols all differ, and large labelled neurodegeneration cohorts are scarce outside a handful of reference centres. Initiatives such as SPIN address biomarker discovery and validation by assembling multicentre observational data spanning genetics, imaging, fluid markers, and structured cognitive assessments [12]; subsets incorporating spontaneous speech tasks tied to standard batteries (for example picture-description prompts compatible with the Western Aphasia Battery [6]) enable language-specific baselines that complement English challenge corpora.

Spanish-focused studies show that language-aware descriptors can still discriminate PPA variants well [26]. Bilingual analyses add a useful caveat: acoustic models often hold up when a patient speaks in a non-dominant language, whereas lexical models grow brittle [27], a contrast worth keeping in mind for multimodal pipelines deployed in multilingual clinics.

2.5 Automatic differentiation of PPA variants from connected speech

Because syndromic overlap complicates visual inspection alone, quantitative speech analyses seek markers that track breakdown profiles hypothesised by neurolinguistic models: reduced speech rate and grammatical complexity in *nfvPPA*; lexical retrieval failures and semantic thinning in *svPPA*; hesitation and phonological loop limitations with spared grammar in *lvPPA* [2].

Empirical studies highlight pause structure, filler usage, lexical diversity, and discourse coherence metrics as discriminative at group level [28], while articulatory tempo measures help separate non-fluent from fluent variants [29, 30]. At the intersection with amyloid biology, connected speech features associate with cortical amyloid burden within PPA cohorts [18], reinforcing the notion that speech captures not only syndrome labels but correlates of underlying pathology, similar in spirit to speech-only screening models targeting amyloid positivity [16, 17].

More recent proof-of-concept work pushes model complexity further: large language models paired with curated linguistic measurements reach high accuracy in retrospective cohorts [31], and architectures trained jointly on acoustic and linguistic embeddings rival expert benchmarks on the multi-class task [32]. These results carry a caveat, though. They stay sensitive to sample size and label noise, and validation across accents, tasks, and cohorts is still an open problem [33].

2.6 Multimodal fusion: pooling, cross-attention, and temporal alignment

Multimodal classifiers must reconcile wildly different sampling rates: acoustic frames at millisecond resolution versus word or sub-word tokens at comparatively coarse timestamps derived from automatic speech recognition [11]. Three broad paradigms describe how audio and text modalities are combined [11]. In early fusion, modality embeddings are pooled along the temporal axis and concatenated into a single fixed-length vector before classification. In late fusion, modality-specific classifiers are trained independently and their outputs are combined at the decision level. Intermediate fusion, the paradigm most relevant to this thesis, processes both modalities jointly through sequence-to-sequence layers that operate on the combined representation before the classification head [13].




Figure 3: Overview of the three multimodal fusion paradigms: early, intermediate, and late.

2.6.1 Self-attention fusion

Self-attention fusion concatenates embeddings from both modalities into a single sequence and processes them with a transformer self-attention layer [34], allowing every position

to attend to every other regardless of modality origin. Multi-head variants distribute this computation across parallel subspaces, each potentially capturing different cross-modal interactions. Adding explicit positional encodings (sinusoidal or learned) to the concatenated sequence before applying self-attention provides ordering information that can help preserve temporal structure and modality boundaries within the joint representation.

2.6.2 Cross-attention and gated variants

Cross-attention introduces directionality: one modality supplies the queries while the other provides keys and values, producing representations of the query modality enriched with information from the other. Gated cross-attention extends this by adding a learned gate, typically a sigmoid-activated linear projection, that scales the cross-attended output element-wise before combining it with the original representation; the gate allows the network to suppress uninformative cross-modal contributions on a per-dimension basis. Bidirectional cross-attention runs two cross-attention operations simultaneously (audio queries text, text queries audio) and merges both enriched representations, so that information flows in both directions.

2.6.3 Selective state-space models: Mamba

An alternative to attention-based fusion comes from the family of selective state-space models. The Mamba architecture [35] processes the concatenated audio-text sequence through a state-space layer that maintains a compressed hidden state updated recurrently along the time axis. A learned, input-dependent selection mechanism controls how much each position influences the hidden state, enabling the model to retain salient information and discard noise dynamically. Because the state update is recurrent rather than pairwise, computational cost grows linearly with sequence length instead of quadratically as in full self-attention, a property that becomes practically relevant when processing long recordings of spontaneous speech.

2.6.4 Cross-modal alignment

Orthogonal to the choice of fusion mechanism, *cross-modal alignment* addresses the problem of establishing temporal correspondence between streams that operate at fundamentally different resolutions. Acoustic encoders such as Wav2Vec 2.0 produce one contextual frame approximately every 20 milliseconds, whereas text encoders emit one embedding per subword token; a single spoken word may span dozens of acoustic frames while simultaneously being split into multiple subword units by the tokeniser. Without explicit alignment, the fusion layer must learn this correspondence from limited labelled data, risking wasted capacity and spurious associations between acoustically and textually distant positions [14]. The CogniAlign architecture [14] addresses this with a word-to-token alignment strategy. Word-level timestamps are first obtained from the automatic speech recognition step (Whisper), determining which acoustic frames correspond to each spoken word. These word boundaries are then mapped to the subword tokenisation of the text encoder, pairing each textual token with the acoustic frames of the word from which it derives. The aligned representations are subsequently processed by gated cross-attention,

where temporal synchronisation encourages the network to attend to co-occurring acoustic evidence (such as jitter, spectral tilt, or prolonged closures) alongside the corresponding lexical content, rather than to correlates from unrelated temporal positions.

This thesis builds on the CogniAlign alignment idea but generalises it along two independent directions. First, where CogniAlign aligns at the word level and attaches every subword token of a word to the same pooled acoustic span, the present work refines the correspondence to the subword granularity itself, introducing two token-to-token schemes: a *uniform* scheme that divides a word’s acoustic frames evenly across its subword tokens, and a *proportional* scheme that allocates them in proportion to each token’s character length. Second, whereas CogniAlign couples its alignment to a single fusion mechanism (gated cross-attention), this thesis decouples the two and treats fusion as a separate axis spanning the seven families surveyed above, of which gated cross-attention is only one; bidirectional cross-attention is adopted as the default in the resulting grid. Both extensions, and their interaction with Wav2Vec layer pooling, are developed and evaluated in the methodology chapter (Section 3).




Figure 4: Cross-modal alignment between acoustic frames and text tokens, comparing unsynchronised input (top) with word-to-token aligned input (bottom).

Fusion architectures developed for speaker and emotion modelling have been shown to transfer across paralinguistic tasks once domain regularisation is added [36]. That lineage is the reason this thesis is willing to port encoder stacks originally tuned on emotion corpora to clinical speech [13].

2.7 Handcrafted acoustic descriptors alongside learned embeddings

Despite the strength of self-supervised waveform encoders, clinically interpretable descriptors remain attractive because they tie predictions to measurable articulatory and phonatory parameters. These descriptors can be grouped into five broad categories: temporal and pause features (pause histograms, speech rate), prosodic contours (fundamental-frequency variation, intensity envelopes), voice-quality measures (jitter, shimmer, harmonics-to-noise ratio), spectral descriptors (formant tracks, spectral moments), and cepstral coefficients (mel-frequency cepstral coefficients). Pause-oriented statistics repeatedly emerge as salient in cognitive decline [37, 38]; voice-quality perturbations may track neurogenic speech motor control deficits [5].

Open-source dementia-speech pipelines demonstrate competitive accuracy when coupling classical descriptors with modern backbones [25], supporting hybrid designs where auxiliary vectors augment attention pooling or concatenate immediately before classification [13]. Whether explicit features complement latent embeddings in practice depends on architecture and task: pathology-linked cues sometimes persist outside generic waveform embeddings, whereas architectures with explicit alignment may subsume part of the prosodic signal [14].

2.8 Synthesis and positioning

Taken together, the literature points to four design commitments reflected in the methodological arc of this thesis [13]: *(i)* evaluate pathology-aware models where biology-grounded labels exist (for example $A\beta$ status alongside syndrome diagnoses), not solely dementia-vs-control proxies; *(ii)* benchmark competing multimodal fusion philosophies using harmonised preprocessing on both an English reference corpus (ADRESSo) and a clinically curated Spanish-speaking subset aligned with SPIN; *(iii)* extend the word-to-token alignment proposed by CogniAlign [14] with two novel token-to-token schemes (uniform and proportional) and evaluate their interaction with Wav2Vec layer pooling in joint ablations; and *(iv)* quantify when handcrafted acoustic descriptors remain complementary after state-of-the-art self-supervised acoustic encoding and explicit audio-text alignment [11, 25].

The following methodology chapter translates these commitments into concrete architectures, training protocols, and evaluation choices.

3 Methodology and Project Development

3.1 Problem Formulation

The general problem addressed in this thesis is the prediction of a clinical label from a single recording of spontaneous speech elicited by a structured picture-description task. Each sample is a pair (x^a, x^t) , where x^a is the raw acoustic waveform and x^t is the corresponding transcript of the utterance, and the goal is to learn a mapping $f_\theta : (x^a, x^t) \mapsto y$, where $y \in \mathcal{Y}$ is a discrete diagnostic label whose domain depends on the task.

We study three tasks, each defined over a distinct clinical cohort and summarised in Table 1. The first is Alzheimer’s disease classification, a binary discrimination between AD patients and cognitively normal (CN) controls on the English ADReSSo cohort [10], with $\mathcal{Y} = \{\text{AD}, \text{CN}\}$; this is the only task with an official held-out test partition. The second is amyloid- β status prediction, a binary proxy-biomarker task defined on a Spanish-speaking subset of the SPIN cohort [12] elicited with the Western Aphasia Battery picture-description task [6], with $\mathcal{Y} = \{\text{positive}, \text{negative}\}$. The third is primary progressive aphasia variant classification, a three-class problem distinguishing the non-fluent/agrammatic, semantic, and logopenic variants [2] defined on the same SPIN subset, with $|\mathcal{Y}| = 3$. Both Spanish tasks are drawn from this single SPIN subset and share the same data tree; neither has an official external test split, so for these we report cross-validated performance only and never construct an artificial test partition.

Both modalities are encoded by frozen self-supervised models. The acoustic stream is processed by a Wav2Vec 2.0 encoder [8] (the specific variant depends on the implementation and is detailed in Section 3.3) producing a frame-level sequence $\mathbf{H}^a = (\mathbf{h}_1^a, \dots, \mathbf{h}_{T_a}^a)$, and the textual stream by a frozen transformer text encoder of the BERT family [9], producing a token-level sequence $\mathbf{H}^t = (\mathbf{h}_1^t, \dots, \mathbf{h}_{T_t}^t)$. Because the two sequences live on different time bases ($T_a \neq T_t$), an alignment step is required to place acoustic and textual representations on a common index before they can be fused, after which a sequence-to-sequence fusion module combines the aligned streams into a representation that feeds a shallow classification head.

Rather than treating this pipeline as a single monolithic system optimised end to end, this thesis formulates it as the composition of six largely independent design axes, each of which can be varied while the others are held at controlled defaults. The modality axis m determines whether the classifier consumes audio alone, text alone, their combination, or additionally a set of handcrafted acoustic features. The fusion axis f selects the sequence-to-sequence architecture that combines the aligned acoustic and textual streams, drawn from a catalogue of seven families. The alignment axis a fixes the scheme that maps acoustic frames onto text positions, comprising the word-to-token baseline of CogniAlign [14] and the two token-to-token schemes proposed in this work. The Wav2Vec pooling axis w chooses between the final encoder layer and a learnable weighted sum across layers. The layer-mixing axis ℓ refines that choice: in place of a single global set of layer weights, an adaptive mixer can learn to combine the encoder’s layers per utterance, or even per token, so that the depth of representation is selected dynamically rather than fixed in advance. Finally, the feature-extractor axis e varies the pretrained encoders themselves, substitut-

ing one audio or text backbone for another, to test how far the downstream trends depend on the specific self-supervised models that supply the two streams.

A concrete system instance is thus parameterised by a tuple (m, f, a, w, ℓ, e) , and the central empirical question is to estimate the marginal contribution of each axis, alone and in interaction, on the three tasks, rather than to identify a single best-performing run. The unit of comparison is therefore the *technique*, not the implementation: two independent codebases replicate the fusion and Wav2Vec-pooling trends under a second software stack, and a third, configuration-unified implementation then consolidates the resulting techniques into a single model (Section 3.3). Building on this decomposition, the experimental programme of Section 4 traverses the axes in a staged sequence, and the synthesis of the strongest choice along each axis into a recommended hybrid architecture is taken up in the Discussion.

3.2 Datasets and Preprocessing

The three tasks defined in Section 3.1 are evaluated on two clinical cohorts, summarised in Table 1: the English ADReSSo benchmark for AD classification, and a curated subset of the Spanish-speaking SPIN cohort that supplies both the amyloid- β and PPA tasks. They differ in language, label space, and the availability of a held-out test partition, which in turn governs the evaluation protocol applied to each (Section 3.7). All recordings originate from structured picture-description tasks, so the input distribution is comparable across cohorts even though the underlying populations and elicitation languages differ. The cohort descriptions and demographic characterisation below follow the companion paper in which we first introduced this SPIN subset [13]; the remainder of the subsection details each cohort and the common preprocessing pipeline that turns raw recordings into the paired acoustic and textual representations consumed by the model.

Table 1: Clinical cohorts, languages, label spaces, and evaluation protocol. Only the English ADReSSo cohort provides an official held-out test partition; the two Spanish tasks are evaluated by stratified cross-validation. Both Spanish tasks are drawn from the same SPIN subset, elicited with the Western Aphasia Battery picture-description (“Picnic Scene”) task.

Task	Cohort	Language	Classes	Evaluation
AD vs. CN	ADReSSo	English	2	Official <code>test-dist</code> (<code>test_acc</code>)
Amyloid- β status	SPIN	Spanish	2	5-fold CV (validation only)
PPA variant	SPIN	Spanish	3	5-fold CV (validation only)

3.2.1 ADReSSo Corpus

The Alzheimer’s disease task uses the ADReSSo-IS2021 corpus introduced for the corresponding Interspeech challenge [10]. The corpus is derived from the English-language DementiaBank Pitt Corpus and comprises 237 spontaneous picture-description recordings, totalling approximately 166 minutes of speech, in which participants describe the Cookie Theft scene from the Boston Diagnostic Aphasia Examination. It is balanced between AD patients and healthy controls and matched for age and sex, which removes

the most obvious demographic confounds and makes it a clean benchmark for isolating speech-based effects. The challenge distributes a fixed train and **test-dist** split, so this is the only task in the thesis with an official external test set and therefore the only one on which we report a held-out test accuracy in addition to cross-validated performance. We use ADReSSo primarily to anchor the architecture against a well-documented English benchmark, as in our companion paper [13], rather than to optimise exhaustively for the challenge leaderboard.

3.2.2 SPIN Clinical Cohort

The two Spanish tasks are defined on a curated subset of the Sant Pau Initiative on Neurodegeneration (SPIN) [12], a longitudinal cohort for biomarker discovery and validation in neurodegenerative disorders maintained at the Hospital de la Santa Creu i Sant Pau in Barcelona. From this cohort we assembled a set of Spanish-speaking participants for whom spontaneous speech recordings, a confirmed amyloid- β status, and, where applicable, a clinical PPA variant diagnosis were all available [13]. Unlike ADReSSo, the elicitation stimulus is the “Picnic Scene” from the Western Aphasia Battery [6], and recordings were captured during routine clinical evaluations, with a variable number of samples per subject ranging from one to seven. This repeated-measures structure is the reason all cross-validation on SPIN is stratified by subject identifier, so that recordings from the same participant never straddle the train and validation folds; reporting at the recording level rather than the subject level would otherwise leak speaker identity and inflate the estimated accuracy.

The subset supports two clinically distinct tasks. The amyloid- β prediction task is a binary problem that asks the model to predict amyloid positivity, a fluid biomarker of AD whose confirmation otherwise requires costly or invasive cerebrospinal-fluid assays or PET imaging. At the subject level it comprises 53 amyloid-positive and 49 amyloid-negative participants, and at the recording level 114 positive and 98 negative samples, leaving the task approximately balanced. The PPA variant classification task is a three-class problem distinguishing the logopenic (lvPPA), non-fluent/agrammatic (nfvPPA), and semantic (svPPA) variants, with 35, 22, and 17 subjects and 104, 63, and 35 recordings respectively. This task is markedly imbalanced (svPPA accounts for fewer than a third of the lvPPA recordings), and the small per-class subject counts make fold-level performance estimates high-variance, so the PPA results are interpreted with corresponding caution.

Tables 2 and 3 report the demographic composition of the two tasks. No statistically significant differences in age, sex, or years of education were found across the groups of either task, assessed with the Wilcoxon rank-sum test for continuous variables and Pearson’s chi-squared test for categorical ones [13]; this matching is important because it makes it less likely that a classifier exploits demographic shortcuts rather than disease-related speech characteristics.

Table 2: Demographic composition of the SPIN amyloid- β task by subject. Values are mean $\pm$ SD or n (%).

	Negative ($N = 49$)	Positive ($N = 53$)
Age (years)	72.8 ± 6.1	72.6 ± 6.8
Female	25 (51.0%)	22 (41.5%)
Male	24 (49.0%)	31 (58.5%)
Education (years)	12.8 ± 6.6	11.8 ± 5.7

Table 3: Demographic composition of the SPIN PPA task by subject. Values are mean $\pm$ SD or n (%).

	lvPPA ($N = 35$)	nfvPPA ($N = 22$)	svPPA ($N = 17$)
Age (years)	72.6 ± 6.5	74.7 ± 4.9	73.4 ± 5.8
Female	15 (42.9%)	9 (40.9%)	11 (64.7%)
Male	20 (57.1%)	13 (59.1%)	6 (35.3%)
Education (years)	10.4 ± 6.4	12.6 ± 5.9	9.8 ± 6.9

3.2.3 ASR Pipeline and Transcript Quality

The textual modality is derived from the audio rather than from gold manual transcripts, so that the pipeline reflects a realistic deployment in which only recordings are available. Each recording is first truncated to its initial forty seconds of speech, which approximates the effective input length of the alignment-based model and, on the SPIN data, avoids capturing the clinician speech that occasionally appears toward the end of an evaluation. The truncated audio is then transcribed automatically with Whisper, though the two stacks differ in configuration: the CogniAligned pipeline runs the Whisper **turbo** model offline to obtain transcripts together with the word-level timestamps its alignment step requires, whereas the SER pipeline reuses precomputed transcripts where they are available and falls back to the lightweight Whisper **tiny** model otherwise. The resulting transcript is tokenised by the text encoder before being aligned with the acoustic stream. Because recognition errors propagate into the textual representation, transcript quality is a genuine source of variance, particularly for the Spanish cohort and for the disfluent or aphasic speech characteristic of PPA; the experiments therefore treat the text modality as one ablatable input among several rather than as a noise-free signal.

3.3 System Architecture

The model follows a two-stream design in which audio and text are embedded independently by frozen self-supervised encoders, aligned onto a common sequence index, combined by a sequence-to-sequence fusion module, and finally pooled and mapped to a task label by a lightweight head (Figure 5). Each stage corresponds to one of the design axes introduced in Section 3.1: the encoders and their pooling fix the acoustic and textual representations, the alignment step determines how the two time bases are reconciled, and the fusion module determines how the aligned streams interact. The ablation programme

depends on this separation: a single axis can be changed while the rest of the pipeline stays fixed.

Both encoders are kept frozen throughout. This is a modelling choice, not a computational shortcut: freezing holds the trainable parameter count in the low millions, concentrates learning capacity in the fusion module under study, and reduces the risk of overfitting the small clinical cohorts. It has a practical payoff as well. Once the encoder outputs are deterministic, every embedding can be computed once and cached, and that cache is what makes a multi-axis grid of experiments tractable at all.

The architecture is realised across three complementary implementations. CogniAligned and SER are the two *primitive* stacks in which the individual techniques were developed and ablated: CogniAligned is the full pipeline and the only stack that implements sub-word token alignment, whereas SER operates on segment-level representations and serves to check whether the fusion and Wav2Vec-pooling trends reproduce under a second, independent implementation. The two stacks differ in their encoders and training schedule, so results are compared across them as directional agreement, not exact reproduction—a caveat that applies wherever cross-stack figures are quoted in what follows. Building on both, `res_model` is a third, configuration-driven implementation that consolidates these techniques into a single model: any combination of modality, fusion, alignment, Wav2Vec pooling, layer mixing, and feature extractor is selected from one configuration, so the strongest choices identified separately on the primitive stacks can be assembled and trained together in a controlled way. The primitive stacks carry the main ablation narrative; `res_model` folds their findings into the unified model from which the recommended architecture is drawn (detailed below). That model is competitive with the primitive stacks across the three tasks without being uniformly the best on any single one. Its value lies in integration and deployment, not in a higher accuracy ceiling.

3.3.1 Acoustic Encoder: Choice and Configuration

Audio is encoded by a frozen Wav2Vec 2.0 model [8], but the two stacks use different variants by design. CogniAligned uses the base configuration (`facebook/wav2vec2-base-960h`): a convolutional feature encoder maps the raw 16 kHz waveform to latent frames at roughly 20 ms resolution (about 50 frames per second), which a twelve-layer transformer context network then contextualises into 768-dimensional hidden states. This model has on the order of 95 million parameters and was trained on 960 hours of English read speech (LibriSpeech); it is therefore monolingual, and CogniAligned applies it unchanged to the Spanish cohorts as well. SER instead uses the cross-lingual XLSR-300M variant, a deeper encoder of roughly 300 million parameters whose twenty-four transformer layers produce 1024-dimensional states and which was pretrained on about 436,000 hours of unlabelled speech spanning 128 languages. The two audio encoders thus differ in depth, width, capacity, and language coverage—one more axis on which the two stacks are not matched.

3.3.2 Wav2Vec Pooling and Layer Mixing

The encoder is used purely as a frozen feature extractor: its weights are never updated, and the frame-level sequence $\mathbf{H}^a$ is read directly from the transformer stack. How that

[Figure to be added]

Figure 5: End-to-end architecture used in this thesis: frozen Wav2Vec 2.0 and BERT-family encoders produce acoustic and textual sequences, which are placed on a common index by the alignment step, combined by the sequence-to-sequence fusion module, and mapped to a task label by a lightweight head.

sequence is formed is the Wav2Vec-pooling axis. The default, used as the controlled setting in every phase except the dedicated pooling study, takes the hidden states of the final transformer layer. The alternative learns one scalar weight per layer, normalises the weights with a softmax, and forms the representation as the weighted sum of all layer outputs (twelve layers for the base encoder in CogniAligned, twenty-four for XLSR-300M in SER). This is motivated by the well-documented observation that the layers of a self-supervised speech encoder specialise in different information (lower layers carry phonetic and speaker detail while higher layers carry more abstract content), so a learnable combination can expose cues that any single fixed layer discards. To further regularise training on the limited clinical data, the SER stack additionally applies on-the-fly acoustic data augmentation, including additive noise, reverberation, and speed perturbation, prior to feature extraction.

The weighted sum just described shares a single set of layer weights across every recording. Making that combination learnable has a clear purpose: instead of committing in advance to the final layer, the model decides for itself which layers to read features from, putting weight wherever the task-relevant signal is strongest. In the recommended base configuration the weighting is a learnable softmax over all twelve transformer layers, and over all twenty-four for the deeper XLSR-300M encoder.

The most informative depth, however, need not be the same for every recording. A refinement evaluated in the unified `res_model` implementation therefore makes the weighting adaptive, replacing the one fixed vector with a learnable *layer mixer* that forms the soft-

max over the layer logits conditionally. It comes in three granularities. A *global* mixer keeps one static weighting for all data. A *sample* mixer predicts a separate weighting per utterance, letting an atypical or noisy recording lean on different layers. A *token* mixer goes finer still, mixing the layers independently at each position. A low temperature ($\tau = 0.1$) sharpens the softmax so the mixer settles on a few layers rather than smearing weight across all of them, and its parameters are trained at about ten times the rate of the rest of the head. The sample mixer is the one carried into the recommended configuration: lightweight, input-conditioned, and in practice biased toward the upper layers.

Mechanically, the mixer acts only on the audio branch and only before fusion, collapsing the layer axis into a single vector per time step; the text embeddings never pass through it. The sample mixer obtains its per-utterance weights by mean-pooling each layer over time and passing the result through a small linear scorer followed by the softmax, so the weighting reflects global properties of the recording; the token mixer instead scores every layer at each position, which amounts to a lightweight attention over layers. This layer mixing should not be confused with attention pooling: the mixer reduces the *depth* (layer) axis before fusion, whereas attention pooling would reduce the *time* axis after fusion, and the recommended configuration uses simple mean pooling for the latter.

3.3.3 Text Encoder: Choice and Configuration

Text is encoded by a frozen transformer from the BERT family [9], which turns the tokenised transcript into a sequence of contextual token-level embeddings $\mathbf{H}^t$. The two implementations again use different members of this family. CogniAligned uses DistilBERT [21], a six-layer encoder distilled from BERT-base that retains most of its language-understanding capability at roughly half the parameters and inference cost, matching the original CogniAlign design. SER uses ModernBERT-base [22], a more recent encoder of about 149 million parameters trained on the order of two trillion tokens, with a native 8192-token context window, rotary positional embeddings, GeGLU activations, and an alternating local-global attention pattern that makes long-context encoding efficient.

As with the acoustic encoder, the text encoder is frozen and its embeddings are precomputed. Holding it fixed within each stack ensures that the text representation is not a confound when fusion, alignment, or pooling is varied. The two stacks do use different text encoders, so this is yet another axis on which they are not matched.

3.3.4 Feature-Extractor Stacks

The choice of pretrained backbone for each stream is itself a design decision rather than a fixed given, and constitutes the feature-extractor axis introduced in Section 3.1. The two primitive stacks already pair different encoders: CogniAligned combines DistilBERT [21] with the English Wav2Vec 2.0 base model, while SER combines ModernBERT [22] with the cross-lingual XLSR-300M. These stacks, however, also differ in their ASR, alignment granularity, and training schedule, so reading the gap between them as an encoder comparison would confound the encoder with everything else that varies. A clean comparison requires changing one encoder at a time while holding the rest of the pipeline fixed.

This is what the unified `res_model` implementation provides through a single-factor extractor ablation. Against a fixed recipe (bidirectional cross-attention fusion, word-to-token alignment, the weighted Wav2Vec stack with the per-utterance sample layer mixer, and word-timestamped CogniAligned transcripts), one component is swapped per run and every other factor is held constant. On the audio side the base encoder is replaced by XLSR-300M, to test whether a deeper cross-lingual model helps the Spanish cohorts; on the text side DistilBERT is replaced by ModernBERT and by the undistilled BERT-base [9], contrasting a lightweight distilled encoder, a recent long-context one, and the original reference. Each variant is run on all three tasks under two random seeds to expose seed sensitivity. The candidates are deliberately drawn from different points of the design space: DistilBERT for its low cost and its role in the original CogniAlign design, ModernBERT for its modern architecture and long context, BERT-base as the undistilled baseline, and the English base versus cross-lingual XLSR audio encoders for the monolingual-versus-multilingual contrast that is most relevant once the same pipeline is applied to Spanish speech. The per-task outcomes, which prove markedly task-dependent, are reported with the experiments.

3.3.5 Cross-Modal Alignment

Before the two streams can be fused they must be placed on a common sequence index, because the acoustic sequence is far longer than the token sequence ($T_a \gg T_t$). CogniAligned resolves this by precomputing an aligned tensor in which each text position is paired with a pooled acoustic representation, and the three alignment schemes differ in how acoustic frames are assigned to text positions. The word-to-token baseline, inherited from CogniAlign [14], uses the word-level timestamps produced by the ASR step to mean-pool the acoustic frames spanned by each spoken word into a single vector, which is then replicated at every subword token of that word. The two schemes proposed in this thesis operate one level finer, at the subword granularity produced by the text tokeniser. The *uniform* scheme takes the frame span of each word and divides it evenly among the word's subword tokens, so that a word split into three tokens contributes three equal slices. The *proportional* scheme instead divides the span in proportion to the character length of each subword token, on the assumption that longer subwords tend to occupy more of the spoken word and should therefore receive proportionally more acoustic frames. Both subword schemes yield a finer-grained audio-text correspondence than the word-level baseline, and the comparison among the three constitutes the alignment axis (Figure 6). SER does not align at the subword level: it fuses segment-level representations and therefore takes no part in the alignment ablation.

3.3.6 Fusion Module

Once the streams are aligned (or, in SER, taken at the segment level), the fusion module combines them into the joint representation that the head consumes. This is the axis the thesis varies most widely. Seven families are implemented and compared, summarised in Table 4, and they fall into three groups (Figures 7–9) that embody three distinct hypotheses about how acoustic and textual evidence ought to interact.

The first group applies **self-attention** to the concatenated audio-text sequence, treating




Figure 6: Audio-text alignment schemes compared in this thesis: the CogniAlign word-to-token baseline, and the two proposed token-to-token schemes (uniform, which spreads each word’s frames evenly across its subword tokens, and proportional, which allocates frames to tokens in proportion to their character length).

both modalities as one undifferentiated set of tokens in which every position may attend to every other. Its premise is that the relevant cross-modal structure will surface on its own once the model is given global context, with no architectural separation between the streams. Three variants probe how much machinery that premise actually needs. Plain self-attention uses a single attention map. Multi-head self-attention splits the computation into several parallel heads, each free to specialise in a different relation, one head might track how prosody aligns with content words while another tracks how pauses fall around syntactic boundaries. The transformer-with-positional-encoding variant goes further and injects sinusoidal position information inside a full transformer block [34], so that the order of frames and tokens is modelled explicitly. That last addition matters for speech in particular, where timing is itself diagnostic and a permutation-invariant model would throw away exactly the rhythm that distinguishes fluent from disfluent production.

The second group, **cross-attention** (Figure 8), keeps the modalities apart and lets one query the other. One stream forms the queries while the other supplies the keys and values, so the fused representation is essentially one modality re-expressed in terms of the most relevant parts of the other. Gated cross-attention extends this with a learned gate that scales the cross-modal contribution, giving the model an explicit means to down-weight the other stream when it is unreliable. That control is not cosmetic here: the text stream is the output of an error-prone ASR step, and a fixed cross-modal coupling would force the model to attend to transcription noise it would be better off ignoring. Bidirectional cross-attention performs the exchange in both directions at once, refining the audio with




Figure 7: Self-attention fusion over the concatenated audio-text sequence.

the text and the text with the audio instead of designating either as the primary stream.

The final family abandons attention altogether. The **Mamba** block [35] is a selective state-space model that mixes the sequence in linear rather than quadratic time, and it carries a very different inductive bias: information is propagated through a recurrent hidden state whose dynamics are themselves input-dependent, which favours smoothly evolving temporal structure over the all-pairs comparison that attention performs. Including it asks whether any gain attributed to fusion is specific to attention or survives under a fundamentally different sequence model.

The same seven families are implemented in both codebases. Fixing the catalogue this way is what lets the fusion axis be read as a property of the technique itself, reproducible across two independent software stacks, rather than an artefact of one implementation.

Table 4: The seven sequence-to-sequence fusion families compared in the experiments, grouped by mechanism.

Group	Fusion family	Core mechanism
Self-attention	Self-attention	Single attention map over the concatenated sequence
	Multi-head self-attention	Parallel attention heads over the concatenated sequence
	Transformer + PE	Transformer block with explicit positional encoding
Cross-attention	Cross-attention	One modality queries the other
	Gated cross-attention	Cross-attention with a learned gate on the cross-modal signal
	Bidirectional cross-attention	Cross-attention exchanged in both directions
State-space	Mamba	Selective state-space mixing in linear time




Figure 8: Cross-attention fusion variants: (a) standard, (b) gated, and (c) bidirectional.

3.3.7 Handcrafted Feature Extraction and Integration

In addition to the learned embeddings, the pipeline can incorporate a set of 55 utterance-level handcrafted acoustic descriptors designed to capture fine-grained paralinguistic cues that are clinically interpretable but not easily recovered by a self-supervised encoder operating on the raw waveform. They are produced by a feature-extraction script adapted from NeuroXVocal [25], which demonstrated the set’s effectiveness on the ADReSSo benchmark, and were extended to the SPIN cohort in our companion paper [13]. Extraction relies on `librosa` together with Praat-based analysis via `parselmouth`; the resulting vector is computed once per recording and stored, so that the identical 55-feature set is reused across every experiment that enables it. The descriptors span the five groups listed in Table 5: temporal and pause statistics, acoustic-prosodic measures, voice-quality measures, spectral descriptors, and cepstral coefficients.




Figure 9: Mamba selective state-space fusion applied to the concatenated audio–text sequence.

Table 5: The five groups of handcrafted acoustic descriptors (55 features in total), extracted with `librosa` and Praat (`parselmouth`).

Group	Descriptors
Temporal & pause	Total speech duration, speech-to-pause ratio, number of pauses, pause-duration statistics (mean, maximum, SD)
Acoustic–prosodic	Pitch (mean, SD, range), intensity, speaking rate, articulation rate
Voice quality	Local jitter, local shimmer, harmonics-to-noise ratio
Spectral	Spectral-centroid statistics, zero-crossing rate, formants F1–F3 (mean, SD)
Cepstral	First 13 MFCCs (mean, SD)

The five groups are grounded in the clinical phenomenology of the target disorders. Pause behaviour is the clearest case: a higher pause rate, longer pauses, and a lower speech-to-pause ratio are well-established markers of Alzheimer’s disease, where they track the lexical-retrieval and speech-planning difficulties that accompany cognitive decline. Voice quality plays a complementary part, with jitter and shimmer registering the subtle phonatory instabilities that can appear before any overt symptom. Such descriptors earn a separate place for two reasons: they are explicit and clinically interpretable, and a self-supervised encoder offers no guarantee of preserving them. Pretrained to recover phonetic and lexical content, the encoder has no objective that obliges it to retain absolute timing or low-level voice-quality statistics, so it may discard precisely the cues a clinician would attend to.

Before the features are used they are normalised, and the normalisation is computed with

care to avoid leakage: each dimension is Z-score standardised using only the mean and standard deviation of the training split of the fold in question, never statistics pooled over the whole dataset.

Integration follows a late-fusion design. The normalised 55-dimensional vector is first projected to the model's 768-dimensional hidden width, then combined with the multimodal stream at one of two late points, prepended as an additional token to the sequence that enters the fusion module, or concatenated to the already-pooled fixed-size embedding immediately before the classification head. Either way the descriptors join only after the self-supervised streams have been formed, so the fusion and cross-attention machinery still operate on the learned representations alone and the handcrafted cues act as an interpretable side-channel rather than something the attention has to disentangle from the waveform.

The descriptors are ablated explicitly in the modality study (Step A), where the audio-plus-text-plus-features cell is compared against the audio-plus-text cell so that any contribution of the handcrafted features is measured against an otherwise identical model. Beyond that controlled comparison, the CogniAligned ablation grids keep the 55-feature vector enabled by default whenever the cohort's feature CSV is available, so the fusion, alignment, and pooling studies are run on top of the audio-plus-text-plus-features input rather than with the features removed.

3.3.8 Classification Heads

The fused sequence is reduced to a single fixed-size vector by a pooling operation (mean pooling in CogniAligned and learned attention pooling in SER), and this vector is passed to a small multi-layer perceptron that outputs the class logits, trained with a cross-entropy objective. The head is kept intentionally shallow so that the bulk of the modelling capacity lies in the frozen encoders and the fusion module under study; this keeps the trainable parameter count low and makes the comparison between fusion families fair, since differences in accuracy can then be attributed to the fusion mechanism rather than to head capacity. The output layer follows the task, with two units for the binary AD and amyloid- β tasks and three for PPA variant classification. All three targets in this thesis are categorical, so the head is always a classifier; the same design would extend directly to a regression target (for instance a continuous amyloid burden rather than a binary status) by replacing the final layer with a single linear unit and the cross-entropy objective with a regression loss, a possibility noted here only to delimit the scope of the present work.

3.4 Unified Implementation: `res_model`

The two primitive stacks established the trends of the individual axes, but each carries incidental design decisions, such as a fixed ASR model, alignment granularity, and training schedule, that make some comparisons impossible to isolate within a single stack. To remove these confounds and to assemble the strongest per-axis choices into one system, the techniques are re-implemented in `res_model`, a configuration-driven trainer in which the modality, fusion family, alignment scheme, Wav2Vec pooling and layer mixer, and the audio and text encoders are all selected from a single configuration. Its results are

reported by technique rather than by stack identity, since a given configuration no longer corresponds to “CogniAligned” or “SER” but to an explicit combination of choices. This unified implementation serves three purposes: it reproduces the main technique trends under one codebase, so that they no longer depend on cross-stack differences; it enables the controlled single-factor studies that the primitive stacks cannot express cleanly, namely the adaptive layer mixer of Section 3.3.2 and the encoder ablation of Section 3.3.4; and it provides the vehicle for the recommended architecture.

That recommended architecture is the configuration the ablations point to, anchored on the English ADReSSo task. It pairs word-timestamped transcripts and a DistilBERT text encoder with the Wav2Vec 2.0 base audio encoder, reads the audio as a weighted stack of layers combined by the per-utterance sample mixer, aligns the streams word-to-token, fuses them with bidirectional cross-attention, and classifies a mean-pooled representation with a shallow head; the handcrafted acoustic features are left off by default and enabled only as an ablation. The bundle is deliberately literature-aligned: its word-to-token alignment follows CogniAlign, while its bidirectional cross-attention fusion is the strongest family from the fusion grid rather than CogniAlign’s own gated cross-attention. It is chosen as the strongest configuration that stays stable across folds and seeds, not the single highest-scoring cell of any one run. As noted above, on the Spanish cohorts some of these choices change: the cross-lingual XLSR-300M audio encoder, in particular, can be preferable to the English base encoder, so the recommended stack is read as task-anchored rather than universal.

3.5 Training Protocol

All models are trained under stratified five-fold cross-validation applied independently to each dataset, so that every fold preserves the class balance of its cohort. The two implementations use schedules tuned to their respective encoders. CogniAligned is optimised with a learning rate of 2×10^{-5} for up to two hundred epochs with early stopping on the validation fold using a patience of fifteen epochs, while SER uses a learning rate of 5×10^{-5} for twenty-five epochs. To separate genuine effects from initialisation noise on these small cohorts, every `res_model` configuration is trained under two fixed random seeds (43 and 3407) and both runs are reported, so that the stability of a result across seeds can be judged alongside its mean. To keep training and inference tractable with frozen encoders, embeddings are precomputed and cached: CogniAligned stores token-aligned tensors as precomputed `.pt` files, and SER maintains an embedding cache. This caching was added to SER after its initial implementation, so all reported inference times correspond to the final cached configuration.

3.6 Ablation Protocol

The experiments isolate the contribution of each design axis by varying one axis at a time while the others are pinned to controlled defaults, so any change in performance can be attributed to the manipulated axis. In the CogniAligned grids those defaults are the final-layer Wav2Vec representation, word-to-token alignment, and bidirectional cross-attention fusion, with audio-plus-text input and the 55 handcrafted acoustic features kept

on whenever the cohort’s feature CSV is available. The only cells without features are the Step A modality cells, which exist precisely to isolate audio, text, or their combination.

The programme runs as a staged sequence. A modality study (Step A) comes first, then a fusion-family comparison (Step B); an alignment study (Tier A) compares the three schemes under a fixed fusion (bidirectional cross-attention or Mamba) at the final Wav2Vec layer; a pooling study (WvLayer) contrasts the final layer with a learnable weighted sum; and a combined study (Align×Wv) crosses the alignment schemes with weighted pooling at bidirectional cross-attention. The strongest per-axis choices are then consolidated and validated in `res_model`, which additionally carries the single-factor layer-mixer (Section 3.3.2) and encoder (Section 3.3.4) studies that the primitive stacks cannot express cleanly.

SER plays a narrower part: it replicates only the fusion (Step B) and Wav2Vec-pooling (WvLayer) trends and has no alignment axis. As noted above, its figures are read as directional support for the CogniAligned trends, not a head-to-head comparison.

3.7 Evaluation Metrics and Baselines

The primary metric is classification accuracy, reported in a manner consistent with each cohort’s evaluation protocol. For ADReSSo we report both the mean cross-validated validation accuracy and the accuracy on the official `test-dist` partition, the latter being the headline figure comparable with the challenge literature. For the two Spanish tasks, which have no official test set, we report validation accuracy only, computed as the mean of the per-fold best validation accuracy across the five folds, and we never substitute an artificially held-out split for the missing test partition. Alongside accuracy, per-sample validation and test inference times are recorded to characterise the compute cost of each technique, since some fusion families are substantially heavier than others. The word-to-token alignment of CogniAlign [14] serves as the literature baseline against which the two proposed token-to-token schemes are measured, and challenge-reported systems on ADReSSo provide the external point of comparison for the AD task. Unless stated otherwise, results are reported as point estimates without significance testing, a limitation revisited in the Discussion.

4 Results

This chapter reports the experimental programme laid out in Section 3. The presentation runs task by task first (Alzheimer’s disease on ADReSSo, amyloid- β status on SPIN, and PPA variant classification on the same SPIN subset), and then opens out into the per-axis ablations that explain where each task’s accuracy comes from.

Two conventions hold throughout. ADReSSo is the only cohort with an official held-out partition, so it is the only task for which a **test-dist** accuracy is quoted alongside cross-validated validation accuracy; the two Spanish tasks are reported on validation folds only, and no artificial test split is ever substituted for the one they lack. Every accuracy is a mean over the five cross-validation folds and the two training seeds (43 and 3407), given as a point estimate. No significance testing was carried out, a limitation revisited in the Discussion.

Unless a table says otherwise, the numbers come from the unified `res_model` implementation run under a single stable recipe (weighted Wav2Vec 2.0 layers, mean pooling, and no handcrafted features) so that one axis moves at a time. Where the older CogniAligned and SER grids supply something the unified runs cannot, namely a second independent stack for the fusion and pooling trends or the legacy alignment sweep on amyloid, those numbers are brought in explicitly and labelled as such.

4.1 Alzheimer’s Disease Classification (ADReSSo)

ADReSSo is the anchor task. It is a balanced, age- and sex-matched English benchmark with a fixed test partition, which makes it the one place in the thesis where a model can be judged against an external, previously unseen split rather than on its own folds.

The recommended configuration pairs bidirectional cross-attention fusion with word-to-token alignment and a weighted stack of Wav2Vec 2.0 layers refined by the per-utterance sample mixer of Section 3.3.2. Averaged over the five folds and both seeds it reaches **88.16%** validation accuracy. Trained once on the full labelled set and evaluated on the official **test-dist** partition, the pipeline scores **83.10%** test accuracy (Table 12), which sits among the stronger systems reported for the challenge and was obtained without any leaderboard-specific tuning.

One result deserves to be flagged early, because it returns in the ablations. On ADReSSo the Mamba fusion block is not merely competitive: it edges the validation grid (88.44% in the legacy Step B sweep) and supplies the 83.10% test figure after full training. That is a genuine finding and it is reported as one. It does not, however, make Mamba the recommended default. On the Spanish PPA task the same block collapses to near chance (Section 4.4), and its ADReSSo advantage is sensitive to the seed. The recommendation therefore keeps the more even-tempered bidirectional cross-attention, trading a little ADReSSo accuracy for behaviour that holds up across all three tasks.

4.2 Amyloid- β Status Prediction (SPIN)

The amyloid task asks the model to recover a fluid biomarker (amyloid positivity) from forty seconds of Spanish picture description, a label otherwise confirmed only by a lumbar puncture or a PET scan. There is no official test partition, so every figure is a five-fold cross-validation accuracy.

Under the unified stable recipe, gated cross-attention is the best plain fusion at 82.88% (Table 8), and the task leans noticeably on the text stream (Table 6). The headline number, though, comes from the configuration this thesis contributes. With character-proportional token-to-token alignment, a final-layer audio representation, bidirectional cross-attention, and the 55-dimensional handcrafted acoustic features switched on, a tuned run reaches **83.90%** validation accuracy (Table 12).

The alignment effect is worth isolating here, because it is the methodological core of the thesis and amyloid is where it shows most clearly. In the legacy Tier A grid (bidirectional cross-attention, handcrafted features, final-layer audio, everything else fixed), character-proportional token-to-token alignment reaches **85.72%**, against **85.13%** for the CogniAlign word-to-token baseline and the same 85.13% for the uniform variant. The proportional scheme is the only one of the three to move the result, and it moves it the right way (+0.59 pp). The unified `res_model` re-implementation reproduces the recipe at 83.90% rather than 85.72%; that gap reflects differences in the training stack, not a reversal of the trend, and the two figures should be read together as reproduction status rather than as competing claims.

4.3 PPA Variant Classification (WAB)

PPA variant classification is the hardest of the three tasks and the one whose numbers carry the most caveats. It is a three-class problem over a small, imbalanced cohort (35, 22, and 17 subjects for the logopenic, non-fluent, and semantic variants), so fold-level accuracy is noisy, and a handful of misassigned recordings shift the mean appreciably.

The best configuration uses bidirectional cross-attention, uniform token-to-token alignment, and a final-layer audio representation, reaching **74.95%** validation accuracy (Table 12). Two patterns set PPA apart from the other tasks. It is the only task that prefers audio to text as a single modality (60.03% against 53.57%, Table 6), which fits the motor-speech and fluency disturbances that mark the non-fluent variant. And it is unforgiving of the wrong fusion: the Mamba block, competitive elsewhere, falls to 52.57% (barely above the majority class), while the cross-attention family stays in the low seventies (Table 8). Given the cohort size, these results are read as indicative rather than definitive.

4.4 Ablation Studies

The per-task numbers above are summaries of a larger grid. This section walks that grid one axis at a time, in the order the pipeline is built: which modalities to feed in, whether handcrafted features help, how to fuse the streams, how to pool the Wav2Vec layers, how to align audio to text, and which pretrained encoders to use. Each axis is varied with the others pinned to the stable recipe, so that a change in accuracy can be charged to the

axis under test. For the two axes that a second, independent implementation also covers (fusion and Wav2Vec pooling), its agreement with the unified runs is noted.

4.4.1 Modality

Naive combination of the two streams is not enough. Concatenating or averaging the audio and text embeddings (Table 6) never beats the better of the two unimodal branches, and on the Spanish tasks it sits well below it. The tasks also disagree about which modality carries the signal: amyloid is text-led (67.85% text against 66.23% audio), whereas PPA is audio-led (60.03% against 53.57%); ADReSSo, like amyloid, prefers text. The practical lesson is that extracting value from both modalities at once needs an attention-based fusion able to learn the correspondence between them. Simple pooling leaves accuracy on the table.

Table 6: Modality ablation per task (`res_model`). SPIN: weighted Wav2Vec 2.0 layers, mean pooling, no acoustic features, per-task default alignment. ADReSSo: minimal-encoder baselines, word→token, final layer. Mean CV validation accuracy (%) over seeds.

Input configuration	Amyloid	PPA	ADReSSo
Audio only	66.23	60.03	53.97
Text only	67.85	53.57	67.72
Audio + Text (mean)	64.12	54.10	65.94
Audio + Text (concat)	64.12	54.10	65.94

Unimodal branches dominate naive fusion; amyloid leans on text, PPA on audio. Attention fusion (Table 8) is required to exploit both modalities.

4.4.2 Handcrafted Acoustic Features

The 55 librosa/Praat descriptors are an optional side-channel, and the ablation (Table 7) shows their contribution to be real but task-dependent. Concatenated after pooling, they lift amyloid by 2.1 pp (78.60% to 80.68%) and ADReSSo by 0.6 pp, but cost PPA 2.3 pp. There is no universal gain. That is why the recommended default leaves the features off and turns them on only where a task earns it, amyloid being the clear case, consistent with the pause- and voice-quality phenomenology of cognitive decline.

4.4.3 Fusion

Fusion is the axis with the widest spread (Table 8, shown graphically in Figure 10), and the three groups behave exactly as their inductive biases would predict. Pooling and plain self-attention stay near chance on the Spanish tasks; self-attention over an undifferentiated audio-text bag has no way to find the cross-modal structure on its own. Positional information and multiple heads help, but the decisive jump comes from keeping the modalities apart and letting one attend to the other. The cross-attention family (plain, gated, and bidirectional) tops every task: between 81 and 83% on amyloid, the low seventies on PPA,

Table 7: Handcrafted acoustic-feature (AF) ablation per task (`res_model`, bidirectional cross-attention, weighted Wav2Vec 2.0 layers). 55-dim librosa/Praat descriptors. *None* = no AF; *token* = AF as an extra sequence token; *concat* = AF concatenated after pooling. Mean CV validation accuracy (%) over seeds. Per-task alignment: amyloid word→token, PPA tok→tok (uniform), ADReSSo word→token.

Acoustic features	Amyloid	PPA	ADReSSo
None	78.60	72.52	86.94
Token integration	n/a*	n/a*	86.94
Concatenation	80.68	70.25	87.54

*Token-level AF integration was only swept on ADReSSo. Concatenated AF helps ADReSSo (+0.6 pp) and amyloid (+2.1 pp) but slightly hurts PPA (−2.3 pp): acoustic features are beneficial but task-dependent, not universal.

and the high eighties on ADReSSo. Mamba is the telling outlier. It is strong on English and amyloid (85.40% on ADReSSo, 77.49% on amyloid) yet collapses to 52.57% on PPA, and that collapse is what rules it out as a general-purpose default despite its ADReSSo standing.

This ranking is the main trend the second, independent SER stack was built to check, and it survives the change of implementation. In the legacy Step B grid the cross-attention and Mamba families again lead within SER (bidirectional cross-attention reaches 85.21% validation on ADReSSo, for instance), even though SER uses different encoders, a different ASR path, and a shorter training schedule. The absolute numbers diverge between the two stacks, which are not a controlled head-to-head, but the ordering of fusion families holds, and reproducing that ordering twice is the whole point of running it twice. The stacks also differ starkly in inference cost. Because CogniAligned reads precomputed token-aligned tensors while the SER timings include the encoders in the loop, per-sample inference differs by more than two orders of magnitude: sub-millisecond against a couple of hundred milliseconds (Figure 11). That gap is a property of the embedding cache, not of the fusion mechanism, and it is what makes a multi-axis grid affordable in the first place.

4.4.4 Wav2Vec Layer Pooling and the Layer Mixer

Two questions sit on the audio branch: whether to read the final Wav2Vec layer or a learnable mixture of all twelve, and, if a mixture, whether a single global set of weights is enough. The pooling table (Table 9) gives a muted answer. A weighted layer stack helps ADReSSo slightly (86.64% against 86.32%), but on the Spanish tasks the final layer is as good or better, clearly so on PPA (74.95% against 72.43%). Depth, once a competent fusion is in place, buys little.

The adaptive layer mixer, evaluated only in the unified implementation, sharpens the picture. Replacing the single global weight vector with a per-utterance *sample* mixer at temperature 0.1 lifts ADReSSo from 87.23% (global mixer, uniform initialisation) to **88.16%**, and the learned weights are interpretable: the mixer concentrates on the top layer (layer 11, weight ≈ 0.20 against the uniform 0.083) while keeping secondary mass

[Figure to be added]

Figure 10: Validation accuracy of the fusion families on the three tasks (values from Table 8). The cross-attention family leads on every task, while Mamba’s drop to chance on PPA stands out as the clearest outlier.

[Figure to be added]

Figure 11: Validation accuracy against per-sample inference time for the fusion families (Step B, ADReSSo). Cross-attention pairs high accuracy with low cost in the precomputed-embedding pipeline; absolute timings depend on whether the encoders are cached or run in the loop.

Table 8: Fusion-technique ablation per task (`res_model`, weighted Wav2Vec 2.0 layers, no acoustic features). Mean CV validation accuracy (%) over seeds. Fixed per-task alignment: ADReSSo tok→tok (char); amyloid word→token; PPA tok→tok (uniform).

Fusion technique	Amyloid	PPA	ADReSSo
Concatenation	64.12	54.10	66.83
Element-wise mean	64.12	54.10	66.83
Self-attention	61.26	54.09	68.65
Transformer (positional)	78.05	59.04	80.58
Multi-head self-attention	75.85	66.66	84.22
Cross-attention	81.00	73.67	86.34
Gated cross-attention	82.88	73.46	86.34
Bidirectional cross-attention	81.64	72.43	87.23
Mamba	77.49	52.57	85.40

Cross-attention variants dominate on every task; Mamba is competitive on English and amyloid but collapses on PPA. Naive concat/mean/self-attention fusions stay near chance on SPIN.

on the lowest layer, and it varies those weights from one utterance to the next instead of settling on a single profile (Figure 12). The finer *token* mixer does not improve on this, so the per-utterance sample mixer is the variant carried into the recommendation: cheap, input-conditioned, and biased toward the upper layers where lexical content concentrates.

Table 9: Wav2Vec 2.0 layer-pooling ablation per task (`res_model`, bidirectional cross-attention). *Final* = last layer; *weighted* = learnable 12-layer mixture. Mean CV validation accuracy (%) over seeds. Per-task alignment: amyloid tok→tok (char), PPA tok→tok (uniform), ADReSSo word→token.

Pooling	Amyloid	PPA	ADReSSo
Final layer	80.75	74.95	86.32
Weighted layers	80.75	72.43	86.64

Weighted layers help ADReSSo slightly; on SPIN the final layer is as good or better, especially on PPA.

4.4.5 Audio–Text Alignment

Alignment is the thesis’s own contribution, so it is examined under two lenses. Under the unified stable recipe (Table 10) the three schemes finish close together on all three tasks: character-proportional token-to-token is best on PPA (73.48%), word-to-token edges amyloid (81.64%), and on ADReSSo the two token-to-token schemes and the baseline are within a few tenths of a point. Subword alignment is therefore not a free lunch under every configuration, and that is worth stating plainly.

The larger effect appears in the configuration where it should: the legacy Tier A grid on amyloid, with handcrafted features and a final-layer audio representation. There, character-proportional alignment reaches 85.72% against 85.13% for word-to-token, with




Figure 12: Learned per-layer weights of the per-utterance sample mixer on ADRessSo ($\tau = 0.1$, twelve Wav2Vec 2.0 layers). Weight concentrates on the uppermost layer with secondary mass on the lowest, and the per-utterance spread is shown as a shaded band around the mean.

the uniform scheme matching the baseline at 85.13%. The proportional allocation, which gives longer subwords proportionally more of the spoken word’s acoustic frames, is the one refinement that helps, and amyloid is where it helps most. Read across both lenses, the honest summary is that token-to-token alignment is a worthwhile, low-cost refinement that pays off on specific task and feature combinations rather than a universal improvement, and that the character-proportional variant is the stronger of the two proposals.

Table 10: Audio–text alignment ablation per task (`res_model`, bidirectional cross-attention, weighted Wav2Vec 2.0 layers). Mean CV validation accuracy (%) over seeds. The token→token schemes are the contribution of this thesis.

Alignment	Source	Amyloid	PPA	ADReSSo
word→token	CogniAlign	81.64	72.43	86.32
tok→tok (uniform)	This thesis	79.79	72.43	86.02
tok→tok (char)	This thesis	80.75	73.48	86.02

Under a stable recipe the three alignments are close on all tasks; character-proportional token→token is best on PPA, while word→token edges amyloid. The larger token→token gain on amyloid appears in the legacy AF + final-layer stack, where the proportional scheme reaches 85.72%.

4.4.6 Feature Extractors

The last axis swaps the pretrained encoders one at a time (Table 11), and the outcome is firmly task-dependent. ADReSSo is best served by the original lightweight pairing of DistilBERT text with the English Wav2Vec 2.0 base encoder (86.12%), and degrades when either is replaced, most sharply when the audio encoder switches to the multilingual XLSR-300M (80.37%). The Spanish tasks pull the other way: XLSR-300M audio gives amyloid its best encoder result (81.40%), and ModernBERT text helps PPA (71.70%). The larger cross-lingual audio model earns its place precisely where the language no longer matches the English pretraining, which is the contrast the ablation was designed to expose, and a reminder that the most capable encoder is not automatically the right one once the task is fixed.

Table 11: Feature-extractor (encoder) single-factor ablation per task (`res_model`; fixed recipe: bidirectional cross-attention, word→token, weighted Wav2Vec 2.0 + sample mixer). One component swapped per run; mean CV validation accuracy (%) over seeds 43 and 3407.

Variant	Text encoder	Audio encoder	Amyloid	PPA	ADReSSo
Baseline	DistilBERT	Wav2Vec2 base-960h	79.84	69.94	86.12
Audio → XLSR	DistilBERT	Wav2Vec2 XLSR-300M	81.40	71.68	80.37
Text → ModernBERT	ModernBERT	Wav2Vec2 base-960h	81.36	71.70	77.38
Text → BERT	BERT-base	Wav2Vec2 base-960h	79.88	69.93	81.91

ADReSSo favours DistilBERT + base-960h; on SPIN, XLSR-300M audio (amyloid) and ModernBERT text (PPA) slightly exceed the baseline.

4.4.7 Unified-Model Confirmation

Taken as a whole, the unified `res_model` runs reproduce the main trends established on the two primitive stacks (the cross-attention family wins fusion, depth beyond the final layer adds little, and subword alignment helps narrowly and task-specifically), while adding the two studies the primitive stacks could not express cleanly: the adaptive layer mixer and the single-factor encoder swaps. Table 12 collects the strongest bundle found for each task. No two of them share the same recipe, which is the thesis’s central empirical message in miniature: the best choice along one axis depends on the task, and a single configuration that wins everywhere does not exist in this grid. What the evidence supports instead is a sensible default (bidirectional cross-attention, word-to-token alignment, a weighted audio stack with the sample mixer, and handcrafted features off) that stays close to each per-task best while remaining stable across folds, seeds, and tasks.

4.5 Error Analysis

A detailed error analysis, including per-class confusion matrices for the three-class PPA task and an inspection of misclassified ADReSSo and amyloid cases, is left as future work. The present study reports point-estimate accuracies without confusion-level breakdowns

Table 12: Best technique bundle per task in the unified `res_model` implementation. SPIN (amyloid, PPA) are Spanish, 5-fold CV validation only; ADReSSo (English) additionally reports official test-dist accuracy after full training. Mean over seeds 43 and 3407.

Task	Best technique stack	Val. acc. (%)	Test acc. (%)
Amyloid (SPIN)	bicross, tok→tok (char), final W_v , +AF	83.90^a	n/a
PPA (SPIN)	bicross, tok→tok (uniform), final W_v	74.95	n/a
ADReSSo (ref.)	bicross, word→token, weighted W_v + sample mixer	88.16	83.10 ^b

^aBest single run (seed 43, lr 1.5×10^{-5} , dropout 0.25) with the 55-d handcrafted acoustic features; the plain-recipe gated cross-attention reaches 82.88%. ^bOfficial ADReSSo test-dist under the full-train protocol (Mamba, word→token, final layer), the recipe that topped the ADReSSo fusion grid.

or significance testing; characterising systematic failure modes, for example confusion between the logopenic and non-fluent PPA variants or the influence of ASR transcription errors on the text stream, would require the per-sample prediction logs and is identified here as a direction for subsequent analysis rather than reported with fabricated figures.

5 Budget

6 Environment Impact (Optional)

7 Conclusions and Future Development

References

- [1] John A. Hardy and Gerald A. Higgins. Alzheimer’s disease: The amyloid cascade hypothesis. *Science*, 256(5054):184–185, 1992.
- [2] Maria Luisa Gorno-Tempini, Argye E. Hillis, Sandra Weintraub, et al. Classification of primary progressive aphasia and its variants. *Neurology*, 76(11):1006–1014, 2011.
- [3] Maria Luisa Gorno-Tempini and Miguel Santos-Santos. Primary progressive aphasia. In *Handbook of Clinical Neurology*, pages 653–663. Elsevier, 2015.
- [4] Fanny M. Elahi and Bruce L. Miller. A clinicopathological approach to the diagnosis of dementia. *Nature Reviews Neurology*, 13(8):457–476, 2017.
- [5] Nicholas Rezaii, Bryan Wong, Paul Aisen, et al. Voiceprints of cognitive impairment: analyzing digital voice for early detection of Alzheimer’s and related dementias. *npj Dementia*, 1(1):35, 2025.
- [6] Andrew Kertesz. *Western Aphasia Battery Test Manual*. Psychological Corporation, San Antonio, TX, 1982.
- [7] Kathleen C. Fraser, Jed A. Meltzer, and Frank Rudzicz. Linguistic features identify Alzheimer’s disease in narrative speech. *Journal of Alzheimer’s Disease*, 49(2):407–422, 2016.
- [8] Alexei Baeovski, Yuhao Zhou, Abdelrahman Mohamed, and Michael Auli. wav2vec 2.0: A framework for self-supervised learning of speech representations. *Advances in Neural Information Processing Systems*, 33:12449–12460, 2020.
- [9] Jacob Devlin, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. BERT: Pre-training of deep bidirectional transformers for language understanding. In *Proc. NAACL-HLT*, pages 4171–4186, 2019.
- [10] Saturnino Luz, Fiona Haider, Sofia de la Fuente, Davida Fromm, and Brian MacWhinney. Detecting cognitive decline using speech only: The ADReSSo challenge. In *Proc. Interspeech*, pages 3780–3784, 2021.
- [11] Kevin Ding, Madhu Chetty, Asal Noori Hoshyar, Tanmoy Bhattacharya, and Britt Klein. Speech based detection of Alzheimer’s disease: a survey of AI techniques, datasets and challenges. *Artificial Intelligence Review*, 57(12):325, 2024.
- [12] Daniel Alcolea, Jordi Clarimón, Maria Carmona-Iragui, et al. The Sant Pau Initiative on Neurodegeneration (SPIN) cohort: A data set for biomarker discovery and validation in neurodegenerative disorders. *Alzheimer’s & Dementia: Translational Research & Clinical Interventions*, 5:597–609, 2019.
- [13] R. Esteve, P. Armas, M. Casals-Salvador, M. A. Santos-Santos, A. Bejanin, and J. Hernando. Deep learning based analysis of spontaneous speech for diagnostic classification and biomarker prediction in Alzheimer’s disease and primary progressive aphasia. In *Proc. The Speaker and Language Recognition Workshop (Odyssey 2026)*, 2026. Companion paper associated with this thesis.

-
- [14] Daniel Ortiz-Perez, Maria Benavent-Lledo, Jaime Rodriguez-Juan, Javier Garcia-Rodriguez, and Daniel Tomás. CogniAlign: Word-level multimodal speech alignment with gated cross-attention for Alzheimer’s detection. *Knowledge-Based Systems*, 329:114264, 2025.
 - [15] Matheus Rodrigues Lima, Amanda Capstick, Fatemeh Geranmayeh, et al. Evaluating spoken language as a biomarker for automated screening of cognitive impairment. *Communications Medicine*, 6(1):6, 2025.
 - [16] Emil Fristedt, Caroline Skirrow, Maria Meszaros, et al. Leveraging speech and artificial intelligence to screen for early Alzheimer’s disease and amyloid beta positivity. *Brain Communications*, 4(5):fcac231, 2022.
 - [17] Patricia Gabirondo, Miguel García-Martínez, Andrea Pozueta-Cantudo, et al. Speech biomarkers predict amyloid status in cognitively unimpaired adults. *Intelligence-Based Medicine*, 12:100306, 2025.
 - [18] Anis Slegers, Georgia Chafouleas, Marie Montembeault, et al. Connected speech markers of amyloid burden in primary progressive aphasia. *Cortex*, 145:160–168, 2021.
 - [19] Wei-Ning Hsu, Benjamin Bolte, Yao-Hung Hubert Tsai, Kushal Lakhota, Ruslan Salakhutdinov, and Abdelrahman Mohamed. HuBERT: Self-supervised speech representation learning by masked prediction of hidden units. *IEEE/ACM Transactions on Audio, Speech, and Language Processing*, 29:3451–3460, 2021.
 - [20] Sanyuan Chen, Chengyu Wang, Zhengyang Chen, Yu Wu, Zhuo Liu, Zhuo Chen, Jinyu Li, Naoyuki Kanda, Takuya Yoshioka, Xiong Xiao, et al. WavLM: Large-scale self-supervised pre-training for full stack speech processing. *IEEE Journal of Selected Topics in Signal Processing*, 16(6):1505–1518, 2022.
 - [21] Victor Sanh, Lysandre Debut, Julien Chaumond, and Thomas Wolf. DistilBERT, a distilled version of BERT: smaller, faster, cheaper and lighter. *arXiv preprint arXiv:1910.01108*, 2019.
 - [22] Benjamin Warner, Antoine Chaffin, Benjamin Clavié, et al. Smarter, better, faster, longer: A modern bidirectional encoder for fast, memory efficient, and long context finetuning and inference. In *Proc. ACL*, pages 2526–2547, 2025.
 - [23] Joël Macoir, Monica Lavoie, Galloup Duboisdindien, and Robert Jr Laforce. Rethinking the diagnosis of primary progressive aphasia: current challenges and future directions. *Frontiers in Aging Neuroscience*, 17, 2025.
 - [24] Zarinah S. Syed, Mohammad S. S. Syed, Margaret Lech, and Elena Pirogova. Tackling the ADReSSo challenge 2021: The MUET-RMIT system for Alzheimer’s dementia recognition from spontaneous speech. In *Proc. Interspeech*, pages 3795–3799, 2021.
 - [25] Nikolaos Ntampakis et al. NeuroXVocal: Detection and explanation of Alzheimer’s disease through non-invasive analysis of picture-prompted speech. *arXiv:2502.10108*, 2025.

-
- [26] Jordi A. Matías-Guiu, José Díaz-Álvarez, Fernando Cuetos, et al. Machine learning in the clinical and language characterisation of primary progressive aphasia variants. *Cortex*, 119:312–323, 2019.
- [27] Lokhan Pugalenthi, Nicolás Montagut, Sandra K. Marques-Kiderle, et al. Can a picture description differentiate the nonfluent/agrammatic and logopenic variants of primary progressive aphasia?: Evidence from Catalan-Spanish bilinguals. *Alzheimer's & Dementia*, 20(S10):e094356, 2024.
- [28] Sunghye Cho, Katheryn A. Q. Cousins, Spanaki Shellikeri, et al. Lexical and acoustic speech features relating to Alzheimer disease pathology. *Neurology*, 99(4):e313–e322, 2022.
- [29] Chiara Cordella, Bradford C. Dickerson, Michele Quimby, Yana Yunusova, and Jordan R. Green. Slowed articulation rate is a sensitive diagnostic marker for identifying non-fluent primary progressive aphasia. *Aphasiology*, 31(2):241–260, 2017.
- [30] Katarina L. Haley, Annie Jacks, Jodi Jarrett, Tyler Ray, Kathryn T. Cunningham, Maria Luisa Gorno-Tempini, and Maya L. Henry. Speech metrics and samples that differentiate between nonfluent/agrammatic and logopenic variants of primary progressive aphasia. *Journal of Speech, Language, and Hearing Research*, 64(12):4935–4956, 2021.
- [31] Nicholas Rezaii, David Hochberg, Michele Quimby, et al. Artificial intelligence classifies primary progressive aphasia from connected speech. *Brain*, 147(9):3070–3084, 2024.
- [32] Charalambos Themistocleous, Bronte Ficek, Kimberly Webster, Dirk-Bart den Ouden, Argye E. Hillis, and Kyrana Tsapkini. Automatic subtyping of individuals with primary progressive aphasia. *Journal of Alzheimer's Disease*, 79(3):1185–1194, 2021.
- [33] Naomi Nevler, Sharon Ash, David J. Irwin, Mark Liberman, and Murray Grossman. Validated automatic speech biomarkers in primary progressive aphasia. *Annals of Clinical and Translational Neurology*, 6(1):4–14, 2019.
- [34] Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N. Gomez, Łukasz Kaiser, and Illia Polosukhin. Attention is all you need. In *Advances in Neural Information Processing Systems*, volume 30, pages 5998–6008, 2017.
- [35] Albert Gu and Tri Dao. Mamba: Linear-time sequence modeling with selective state spaces. *arXiv preprint arXiv:2312.00752*, 2023.
- [36] Amir Rahim Naini, Lucas Goncalves, Ali N. Salman, et al. The INTERSPEECH 2025 challenge on speech emotion recognition in naturalistic conditions. In *Proc. Interspeech*, pages 4668–4672, 2025.
- [37] Nelli Imre, Roland Balogh, Gábor Gosztolya, et al. Temporal speech parameters indicate early cognitive decline in elderly patients with type 2 diabetes mellitus. *Alzheimer Disease & Associated Disorders*, 36(2):148–155, 2022.

-
- [38] Tingting Zheng, Xiaohui Wang, Xiaohan Peng, et al. Survey of neurocognitive disorder detection methods based on speech, visual, and virtual reality technologies. *Virtual Reality & Intelligent Hardware*, 6(6):421–472, 2024.

Appendices